

1 IEEE 754 Number Representation

As you can see in your textbook, the IEEE754 Floating Point representation is composed of three parts, the Mantissa Sign, S , the Signed Exponent, E , and the Mantissa Magnitude, M . In single precision floating point representation, the Signed Exponent, E , is 8 bits, whereas the Mantissa Magnitude, M , is composed of the remaining 23 bits. In double precision floating point representation, the Signed Exponent, E is 11 bits, whereas the Mantissa Magnitude, M , is composed of the remaining 52 bits. In both cases, the hidden-1 representation for the Mantissa Magnitude holds, effectively extending its representational power by one bit.

The value of a single precision IEEE754 Floating Point number is typically given by the following formula:

$$N = (-1)^S 2^{E-127} (1.M) \quad (1)$$

Yet, one of the things to keep in mind is that this interpretation only holds for $0 < E < 255$. For $E = 0$ (i.e., E being the bit string “00000000”) and for $E = 255$ (i.e., E being the bit string “11111111”) alternate value interpretations hold as given below.

Condition	N value
$E = 255$ and $M \neq 0$	NaN
$E = 255$ and $M = 0$	$(-1)^S \infty$
$E = 0$ and $M \neq 0$	$(-1)^S 2^{-126} (0.M)$
$E = 0$ and $M = 0$	$(-1)^S 0$

Similarly, the following interpretations hold for the case of *double precision* IEEE754 Floating Point numbers:

Condition	N value
$E = 2047$ and $M \neq 0$	NaN
$E = 2047$ and $M = 0$	$(-1)^S \infty$
$0 < E < 2047$	$(-1)^S 2^{E-1023} (1.M)$
$E = 0$ and $M \neq 0$	$(-1)^S 2^{-1022} (0.M)$
$E = 0$ and $M = 0$	$(-1)^S 0$

Adding and subtraction are the most difficult of the elementary operations for floating-point operands. Here, we deal only with addition, since subtraction can be converted to addition by flipping the sign of the subtrahend. Consider the addition:

$$(\pm s1 \times b^{e1}) + (\pm s2 \times b^{e2}) = \pm s \times b^e \quad (2)$$

Assuming $e1 \geq e2$, we begin by aligning the two operands through right-shifting of the significand $s2$ of the number with the smaller exponent:

$$\pm s2 \times b^{e2} = \frac{\pm s2}{b^{e1-e2}} \times b^{e1} \quad (3)$$

If the exponent base b and the number representation radix r are the same, we simply shift $s2$ to the right by $e1 - e2$ digits. When $b = r^a$ the shift amount, which is computed through direct subtraction of the biased exponents, is multiplied by a . In either case, this step is referred to as alignment shift, or preshift (in contrast to normalization shift or postshift, which is needed when the resulting significand

s is unnormalized). After the alignment shift, the significands of the two operands are added to get the significand of the sum.

When the operand signs are alike, a single-digit normalizing shift is always enough. For example, with IEEE754 format, we have $1 \leq s < 4$, which may have to be reduced by a factor of 2 through a single-bit right shift (and adding 1 to the exponent to compensate). However, when the operands have different signs, the resulting significand may be very close to 0 and left shifting by many positions may be needed for normalization.

Figure 1 shows a floating-point addition example:


E=10001010;	S=1.111000000000000000000000	
+ E=10001000;	S=1.100000000000000000000000	
		
E=10001010;	S=1.111000000000000000000000	
+ E=10001010;	S=0.011000000000000000000000	Alignment shifting
<hr/>		
E=10001010;	S=10.010000000000000000000000	Sum
E=10001011;	S=1.001000000000000000000000	Normalization

Figure 1: floating-point addition

Figure 2 shows a floating-point subtraction example:

E=10001010;	S=1.111000000000000000000000	
- E=10001010;	S=1.110000000000000000000000	
<hr/>		
E=10001010;	S=0.001000000000000000000000	Difference
E=10000111;	S=1.000000000000000000000000	Normalization

Figure 2: floating-point subtraction

Floating-point multiplication is simpler than floating-point addition; it is performed by multiplying the significands and adding the exponents:

$$(\pm s_1 \times b^{e_1}) \times (\pm s_2 \times b^{e_2}) = \pm (s_1 \times s_2) b^{e_1+e_2} \quad (4)$$

Postshifting may be needed, since the product $s_1 \times s_2$ of the two significands can be unnormalized. For example, with the IEEE format, we have $1 \leq s_1 \times s_2 < 4$, leading to the possible need for a single-bit right shift. Also, the computed exponent needs adjustment if a normalization shift is performed.

Figure 3 shows a floating-point multiplication example:

E=10001010;	S=1.100000000000000000000000	
* E=00010001;	S=1.100000000000000000000000	
<hr/>		
E=10011011;	S=10.010000000000000000000000	Product
E=10011100;	S=1.001000000000000000000000	Normalization

Figure 3: floating-point multiplication

Similarly, floating-point division is performed by dividing the significands and subtracting the exponents:

$$\frac{\pm s_1 \times b^{e_1}}{\pm s_2 \times b^{e_2}} = \pm \frac{s_1}{s_2} \times b^{e_1-e_2} \quad (5)$$

The ratio s_1/s_2 of the significands may have to be normalized. With the IEEE754 format, we have $1/2 < s_1/s_2 < 2$ and a single-bit left shift is always adequate. The computed exponent needs adjustment if a normalizing shift is performed.